

AskMoses.AI



System Manual — How Every Number Is Calculated

Reference guide for answering rule questions
and telling intended behaviour apart from real defects

Version 1.0 · July 2026

Based on the `fix/dashboardData` branch

Table of Contents

1. How to Use This Manual
2. Core Concepts
 1. How a call gets scored
 2. The two scales: 0–100 internally, 0–5 on screen
 3. Colour bands and thresholds
 4. Section weight and "critical" flags
 5. Call outcome: Stage 1 and Stage 2
 6. Where calls come from
 7. Organisation scoping and roles
3. The Owner Dashboard
4. Analytics
5. Script Intelligence (Insights)
6. Script Gap Detection
7. History
8. The Rep's Dashboard (My Page)
9. Calls and Call Detail
10. Team Command Center
11. Intent Analysis
12. Marketing Intelligence
13. Admin Panel and Billing
14. Known Defects — Do Not Report These
15. Quick Reference

1. How to Use This Manual

This manual exists so that when someone says "*this number looks wrong*", you can check here first and answer with confidence — is it working as designed, is it a known limitation, or is it genuinely broken?

Every metric in this document is described in three parts: **what it is**, **how it is calculated**, and **why it works that way**. The "why" matters most: almost every number that looks strange at first glance is the result of a deliberate trade-off, and knowing the trade-off is what lets you answer a prospect or a teammate without escalating to engineering.

The three callout boxes

Throughout the manual you will see three kinds of highlighted box. They mean very different things:

WHY THIS EXISTS

The reasoning behind a design decision. If someone asks "why does it do that?", the answer is here.

OFTEN MISTAKEN FOR A BUG

Behaviour that looks like a bug but is intentional. If a report matches one of these, it is not a defect — explain the reasoning instead.

KNOWN DEFECT — ALREADY REPORTED

A genuine defect that has already been identified and verified against real production data. No need to report it again; it is on the fix list. Section 14 collects all of them in one place.

Trust markers on data

Not every number in the product has the same standing. Some are arithmetic on real calls; some are the AI's opinion; some are placeholder values left over from the demo phase. Wherever it matters, you will see one of these tags:

CALCULATED

Arithmetic performed on real call data. Reproducible and auditable.

AI OPINION

A number the language model was asked to produce. It looks like a statistic but no code verifies it. Treat it as a qualified guess, not a measurement.

PLACEHOLDER

A fixed value hardcoded during the demo phase. It does not come from data at all and will not change no matter what the client does.

This distinction is the single most useful thing in this manual. The most common category of "false bug" is someone noticing that an **AI OPINION** number does not reconcile with a **CALCULATED** one. They are not supposed to reconcile — they are different kinds of claim.

2. Core Concepts

2.1 How a call gets scored

Every call, regardless of how it enters the system, follows the same path:

1. **Audio arrives** — either uploaded manually or pulled automatically from the CRM.
2. **Whisper transcribes it** into text.
3. **GPT reads the transcript** against the organisation's *active script* (the rubric).
4. **The AI returns**, in one structured response: a score from 0 to 100 for each section of the rubric, a written justification per section, a list of strengths, a list of improvements, a summary, and its reading of how the call ended.
5. **The overall score is computed** from those section scores.

The rubric's five default sections are Discovery, Problem Agitation, Offer Presentation, Objection Handling, and Close & Next Steps. An organisation can define its own sections through the Script Builder, and the system adapts — nothing is hardcoded to those five names.

The overall score is a simple average

```
overall_score = average of all section scores  
(plain arithmetic mean – every section counts equally)
```

WHY THIS EXISTS

The score measures **execution quality**, not results. A rep can run a textbook call and still lose the deal to a prospect with no budget; another can stumble through a call with a customer who was always going to buy. Keeping the score independent of the outcome is what makes it useful for coaching — it answers "how well did you sell?", not "did you win?". The outcome is tracked separately, as its own field.

OFTEN MISTAKEN FOR A BUG

The score is **not capped by the outcome**. A call that did not close can still score 95, and a call that closed can score 40. This is deliberate and reps notice it. It is not a bug — it is the whole point of separating quality from results.

2.2 The two scales: 0–100 internally, 0–5 on screen

This is the most frequent source of confusion, so it is worth being precise.

The AI is asked to score on a **0 to 100** scale, and 0–100 is what the database stores. Every screen that shows a score divides that number by 20 and displays it with one decimal, on a **0 to 5** scale.

```
displayed score = stored score ÷ 20
```

```
84 / 100 → shown as 4.2 / 5
```

```
72 / 100 → shown as 3.6 / 5
```

WHY THIS EXISTS

The AI produces better, more consistent judgements when it has a wide range to work with — the gap between an 82 and an 87 is a distinction it can reason about, whereas forcing it to pick between 4 and 5 loses that nuance. But reps and owners read a 0–5 rating far more naturally. So the system reasons in 100 and speaks in 5.

The scoring bands the AI is given are: 90–100 textbook execution, 75–89 strong, 60–74 adequate, 40–59 needs work, 0–39 poor or absent.

OFTEN MISTAKEN FOR A BUG

If someone reports "the score in the database doesn't match the screen", check whether they are comparing 0–100 to 0–5. A database value of 84 showing as "4.2" on screen is correct behaviour, not data corruption.

2.3 Colour bands and thresholds

Scores are colour-coded consistently across the product:

Band	Stored (0–100)	Displayed (0–5)	Meaning
Green	85 and above	4.25 and above	Strong execution
Amber	70 to 84	3.5 to 4.2	Adequate, room to improve
Red	below 70	below 3.5	Needs coaching attention

A separate threshold, **95 and above**, marks a call as a "Perfect Call" in the Analytics achievements.

OFTEN MISTAKEN FOR A BUG

One screen uses slightly different cut-offs. On the call detail page, the *written* feedback fallback text switches tiers at 80 and 60, not 85 and 70. This means a call scoring 82 shows an amber bar next to more complimentary text. It is a minor inconsistency that exists because the two systems were built for different purposes, and it is on the cleanup list — but it is not data corruption.

2.4 Section weight and "critical" flags

When an owner builds a script, they can assign each section a **weight** and mark sections as **critical**. Both are saved with every call.

KNOWN DEFECT — ALREADY REPORTED

Neither field currently affects the overall score. The score is a plain average regardless of the weights configured, and marking a section critical does not penalise the call. The fields are captured, stored, and displayed — but never multiplied into the result. An owner who sets "Closing = 50% weight" will see no change in any score. This is verified in the code and is the most consequential gap between what the product appears to offer and what it does.

The `SCHEMA.md` file in the repository describes the overall score as a "weighted average of section scores". That documentation is wrong and predates the current implementation. This manual reflects the actual

behaviour.

2.5 Call outcome: Stage 1 and Stage 2

The product tracks two separate notions of "did this deal happen", and conflating them is a common source of confusion.

Stage 1 — Initial Result

How the *call* ended. One of four values:

Value	Meaning
closed	The deal was closed on this call
not_closed	The call finished with no deal
partial	A follow-up was scheduled, or another decision-maker is involved
no_outcome	No result recorded — including calls that never finished processing

Stage 2 — Actual Close

Whether the client *actually paid*. Marked manually by the owner on the call detail screen: `paying`, `not_paying`, Or `pending`.

WHY THIS EXISTS

A prospect saying "yes, let's do it" on the call and a prospect whose payment clears are two different business events, often weeks apart. Stage 1 is what the AI can detect from the conversation; Stage 2 is ground truth that only a human knows. Keeping them separate is what will eventually allow the system to learn which call behaviours predict real revenue rather than just verbal agreement.

OFTEN MISTAKEN FOR A BUG

Stage 2 can be set, and the marker shows its current value — but **nothing consumes it beyond that**. There is no report comparing predicted intent against who actually paid, and the intent-at-close snapshot captured for that future analysis is never read anywhere. If someone asks "where do I see my Stage 2 numbers?", the honest answer today is "nowhere yet". This is an incomplete feature, not a broken one.

2.6 Where calls come from

Two ingestion paths, both ending in the same place:

- **Manual upload** — the owner uploads an audio file or transcript through the upload screen.
- **CRM webhook** — GoHighLevel notifies the system when a call is recorded, and the pipeline fetches, transcribes, and analyses it automatically. This is the Pro plan feature.

Automatic ingestion can fail for reasons that have nothing to do with selling: no recording available yet, transcription failure, the rep not linked to a CRM user. These calls land in the database with no outcome recorded.

OFTEN MISTAKEN FOR A BUG

Calls that failed to process still count in the Avg Close Rate denominator (see section 3). A batch of failed transcriptions will pull the close rate down even though no sales were lost. If the close rate drops with no sales explanation, pipeline failures are the first thing to check.

2.7 Organisation scoping and roles

Every query is scoped to the active organisation. No screen ever mixes data across clients.

Role	Home after login	Can see	Cannot see
Rep (trainer)	My Page	Own calls and scores, own coaching	Team data, other reps, admin
Owner	Dashboard	Everything for their organisation	The SaaS admin panel, cost/COGS
Admin	Admin Panel	Everything, all organisations	—

Owners never see LLM cost or gross margin — those are filtered out of the billing responses by design, not by UI hiding.

3. The Owner Dashboard

The owner's home screen. Answers: is the team selling well, who needs attention, and what should we fix in the script?

3.1 Avg Close Rate CALCULATED

Avg Close Rate = closed calls ÷ total calls × 100

Scope: the entire organisation, all time.
Every call weighs exactly the same.

WHY THIS EXISTS

This used to be calculated as the average of each rep's individual close rate — which meant a rep with 1 call counted as much as a rep with 50. On real client data that distortion was large: one organisation showed 43% under the old method and 54% under the correct one; another showed 59% when the true figure was 43%. Counting call-by-call is the only way the number reflects the business rather than the shape of the team.

The tooltip on the card states this in plain language for the client. There is no date filter — it is deliberately a lifetime figure, so it is stable and hard to game.

OFTEN MISTAKEN FOR A BUG

The denominator is **every call**, including ones that never got analysed. This keeps the rule simple enough to explain in one sentence, at the cost of letting pipeline failures depress the number. On current client data this costs nothing — almost every organisation has zero calls without an outcome — but it is a known trade-off.

OFTEN MISTAKEN FOR A BUG

The delta underneath the card (" +7 pts since week 1") is **not** measured over the same period as the big number. The headline is all-time; the delta compares the first and last week within a window of at most six weeks and at most 200 calls. They are different questions — "how do we do overall?" and "are we trending up recently?" — shown next to each other without a visual cue.

3.2 Team Avg Call Score CALCULATED

Team Avg Call Score = average of each rep's average score
(only reps who have made at least one call)

WHY THIS EXISTS

Reps with zero calls have a stored score of 0. Including them would drag the team average down every time someone new is invited, making the metric jump for reasons that have nothing to do with performance. Excluding them keeps the number meaningful.

OFTEN MISTAKEN FOR A BUG

This one is an average of averages, unlike the close rate — so a rep with 1 call does influence it as much as a rep with 50. That is inconsistent with how the close rate now works, and is a fair candidate for the same treatment.

3.3 Total Calls and Active Sales People CALCULATED

Total Calls is the sum of every rep's lifetime call count. **Active Sales People** counts reps whose invitation has been accepted — including those who have not yet made a call. Note that these two use different populations, which is why the numbers can look mismatched (five active people, zero calls).

3.4 Correlation Engine MISLEADING LABEL

Presented as "Correlation Engine — What Drives Closes", with High / Medium / Low correlation badges per rubric dimension.

What it actually shows:

score = the team's average score in that dimension

badge = that same score, banded (≥ 85 High, ≥ 70 Med, else Low)

KNOWN DEFECT — ALREADY REPORTED

No correlation is calculated anywhere — not at any call volume. The badges are the team's average score wearing a statistical costume. The panel never compares closed calls against lost calls, which is what a correlation would require. "Objection Handling — High correlation" means only "the team scores above 85 in Objection Handling", not "Objection Handling drives closes".

This matters because it can invert the advice. If the team scores 60 in Objection Handling but nearly every won deal had a high score there, that is the single biggest lever in the business — and the panel would show it in red as "Low correlation", saying the opposite of the truth.

Below three calls, the panel honestly renames itself "Coaching Drivers — Team Averages by Dimension" and shows a disclaimer. Above three calls, the disclaimer disappears and the statistical title takes over. Three calls is far too low a bar for that promise.

WHY THIS EXISTS

A real correlation is feasible — the data needed already exists on every call (per-dimension scores plus the outcome). The measure would be the difference in average score between calls that closed and calls that did not, per dimension. It has not been built yet; the current panel was the interim step.

3.5 Team Health CALCULATED

A per-rep list showing activity recency, close rate, and a trend arrow.

Status	Rule
Active	Last call within 1 day
Recent	Last call within 7 days
Away	No call in over 7 days, or never

The close rate and delta shown here are **not** recalculated live — they are read from a cached column on the rep record, refreshed only when a new call is analysed. The delta compares the last 7 days against the entire history before that.

OFTEN MISTAKEN FOR A BUG

If a rep's stats look stale, this cache is why. It refreshes on call analysis, not on page load. A call that entered through an unusual path (a seed, a manual retry) may not have triggered a refresh, leaving the rep's close rate outdated even though their calls are all in the system.

OFTEN MISTAKEN FOR A BUG

A delta of exactly zero shows an **upward** green arrow, because the rule is "zero or above = up". A rep with no change looks like a rep improving.

3.6 Close Rate Trend CALCULATED

Up to six weeks of team close rate, one point per week, weeks running Monday to Monday.

WHY THIS EXISTS

A week with no calls shows as a **gap in the line, not a zero**. Plotting it as 0% would read as "we closed nothing that week", when the truth is "nobody called that week". The distinction changes what an owner does next, so the chart refuses to guess.

For a rep whose calls all fall within a single week, the chart switches to a per-call cumulative view (C1, C2, C3...) instead of weekly points, because a single weekly point cannot draw a line.

3.7 Score by Sales Person CALCULATED

A grid of reps against rubric dimensions. The highest score in each column is highlighted green; anything below 70 shows red. Based on the 200 most recent calls.

3.8 AI Insights AI OPINION CALCULATED

Four cards mixing real arithmetic with AI-written prose. The numbers inside them are computed in code; the sentences around them are generated. The "Coaching ROI" card quotes the close rate from the same source as the headline card, so those two always agree.

KNOWN DEFECT — ALREADY REPORTED

The "Revenue Leak" card is broken by a scale error. It compares section scores — which are stored 0–100 — against a threshold of 3.5, which was written for the 0–5 scale. Verified against production data: only 4.9% of section scores fall below 3.5, and those are catastrophic scores like 2/100, not mediocre ones like 3.5/5. The result is that the card almost always reads "0 of N sales people score below 3.5", and prints sentences like "Marcus scores 92/5" — a score of 92 out of 5. This is a real defect, already identified.

4. Analytics

A deeper cut of call data. Everything on this page is calculated in the browser from real calls — there is no AI involvement at all.

4.1 Performance Trend CALCULATED

Average call score grouped by day.

4.2 Top Improvement Areas CALCULATED

The average of each rubric section across all loaded calls, sorted worst-first.

4.3 Outcome Metrics CALCULATED

Counts of closed, partial, not closed, and no outcome, plus a close rate using the same rule as the dashboard (closed ÷ total).

4.4 Trainer Conversion Leaderboard CALCULATED

Close rate per rep, ranked. Bars turn green at 50% and above.

OFTEN MISTAKEN FOR A BUG

This leaderboard groups reps by **display name**, not by their unique ID. Two reps with the same name would be merged into one row. Unlikely, but worth knowing.

4.5 Achievements CALCULATED

Badge	Rule
Master Coach	Rep with the highest average score
Perfect Calls	First rep with a call scoring 95 or above
Rising Star	Compares the last 3 calls against the first 3 calls

OFTEN MISTAKEN FOR A BUG

"Rising Star" compares the last three calls of the **whole organisation** against the first three of the whole organisation, then credits the rep who made the most recent call. It can therefore name a rep based on other people's improvement. It is a demo-grade heuristic, not a real measure.

OFTEN MISTAKEN FOR A BUG

This page does not paginate. It loads roughly the 1,000 most recent calls and analyses those. An organisation past that volume will see analytics based on a truncated window, with no warning on screen.

5. Script Intelligence (Insights)

Reads recent transcripts and proposes an improved version of the script. The owner approves or rejects, and approving activates the new script in production.

Sample: up to 7 recent calls with transcripts, each truncated to 1,500 characters.

KNOWN DEFECT — ALREADY REPORTED

Almost every number on this screen is the language model's opinion, presented with the visual authority of a measurement. Specifically: the **Playbook Health Score** (0–100) is asked for directly in the prompt — no code computes or validates it. The **per-section scores** are likewise the AI's read. The **effectiveness label** (good / room to improve / poor) is supposed to follow the health score, but the code never checks that it does — if the model mislabels itself, the screen shows the wrong label. Most seriously, the **uplift figures** on top closer phrases (" +18%") are invented: the prompt asks for a string "e.g. +18%" with no instruction on how to derive it, and no uplift is calculated anywhere in the system.

This matters more than elsewhere because these numbers sit next to buttons that really do swap the organisation's active script. A business decision is being made on the strength of figures that have no statistical basis.

WHY THIS EXISTS

The feature is genuinely useful as *qualitative* analysis — the AI reading seven real transcripts and suggesting sharper language is real value. The problem is presentation: qualitative judgement dressed as quantitative evidence. Reframing these as "AI assessment" rather than scores and percentages would keep the value and drop the false precision.

6. Script Gap Detection

Finds friction points in real conversations that the current script does not address, and proposes targeted rewrites.

Sample: exactly 3 calls with full transcripts. Results are cached for 7 days, then recomputed on next view.

Frequency badge **CALCULATED**

$$\text{frequency} = \text{calls where the gap appears} \div \text{calls analysed} \times 100$$

WHY THIS EXISTS

This is the one AI feature in the product where the percentage is honest. The model identifies *which* calls show the friction, but the percentage is then computed in code from that count — it is not a number the AI made up. Gaps the model reports without pointing to a real call are discarded before saving. This is the pattern the other AI features should follow.

OFTEN MISTAKEN FOR A BUG

With a denominator of 3, the only possible values are 33%, 67%, and 100%. The badge is colour-coded as if it were a fine-grained statistic (red at 60% and above), but it is really reporting "this appeared in 1, 2, or 3 conversations". The percentage is accurate; the precision it implies is not.

The **severity** rating (high / medium / low) is the AI's judgement, validated only against the list of allowed values.

7. History

A searchable table of analysed calls. No aggregate metrics — it displays stored scores and outcomes without recomputing anything.

Calls from the same contact are grouped into a single row with a "View All" action. Paginated at 10 per page in the browser.

8. The Rep's Dashboard (My Page)

The individual seller's view. Deliberately excludes coaching notes written for the owner's eyes.

8.1 The four KPI cards CALCULATED

Score, Close Rate, Calls, and Closed — each with a window selector of 2, 4, or 6 weeks, defaulting to 6.

Within the selected window:

score = weighted by call volume, not a plain average of weeks

close rate = total wins ÷ total calls in the window

WHY THIS EXISTS

Weighting by volume means a week with 12 calls influences the number more than a week with 1. Averaging the weekly percentages instead would let a single-call week swing the figure wildly — one lucky call would read as a 100% week and drag the average up.

OFTEN MISTAKEN FOR A BUG

The delta compares the selected window against **one single week** immediately before it — not against an equivalent window. So "last 6 weeks" is compared to the 7th week alone. That is a noisy comparison and arguably the wrong one, but it is the current design.

OFTEN MISTAKEN FOR A BUG

The small print under each card ("of 47 total") is always the rep's **lifetime** figure, even when the card above shows a 2-week window. Two scopes in one card, by design, to give context.

OFTEN MISTAKEN FOR A BUG

A delta of exactly zero is hidden rather than shown as "0". Reps sometimes read the missing delta as a bug.

8.2 Empty weeks

A week with no calls is flagged as empty rather than scored as zero, and is excluded from comparisons — the same reasoning as the trend chart in section 3.6.

9. Calls and Call Detail

The team's call list and the full analysis of a single call.

The detail screen shows the overall score, per-section bars with the AI's written justification, strengths, improvements, the intent breakdown, the outcome badge, and — for owners only — coaching notes and the Stage 2 marker.

OFTEN MISTAKEN FOR A BUG

A section marked **critical** raises a red alert when it scores 40 or below on this screen. The coaching email uses a different threshold — 60 — so the email flags calls the screen does not. Same concept, two numbers, because they were built separately.

KNOWN DEFECT — ALREADY REPORTED

The "Coaching notes" textarea on the call detail screen is not wired to anything. It has no save action and no persistence — anything typed into it is lost on navigation. It is leftover demo UI.

10. Team Command Center

The owner's per-rep deep dive: behavioural profile, best and worst calls, and AI coaching recommendations.

10.1 Behavioural profile CALCULATED

For each dimension of the active script: the rep's average, the team's average, and the gap between them.

WHY THIS EXISTS

The dimensions come from the organisation's *active script* rather than from each rep's own calls, so that every rep is measured on the same axes and can be compared horizontally. If each rep's chart used their own call history, the charts would not line up.

10.2 Live stats versus cached deltas

OFTEN MISTAKEN FOR A BUG

On this screen the rep's close rate and score are **recalculated live** from their calls, but the deltas next to them still come from the cached column described in section 3.5. So a correct, up-to-date score can sit next to a stale or zeroed delta. The live recalculation was added precisely because the cache is unreliable — but only half the card was migrated.

10.3 Coaching recommendations AI OPINION

The metrics fed *into* the AI prompt are computed in code from the rep's last 20 calls. The recommendations that come back are free text, and the prompt forbids citing specific calls. No fabricated statistics appear in the output.

OFTEN MISTAKEN FOR A BUG

The close rate used inside these recommendations covers the rep's **last 20 calls**, while the dashboard shows their lifetime rate. The same rep can therefore be described as closing at 45% in a recommendation while the dashboard says 38%. Both are correct for their scope.

OFTEN MISTAKEN FOR A BUG

If the AI call fails, the system silently falls back to pre-written recommendations belonging to one of four fictional demo personas. A real rep can be shown generic advice with no indication that it is not personalised.

11. Intent Analysis

Scores how ready each prospect was to buy, independent of how well the rep performed.

11.1 The Intent Index CALCULATED

The AI reads the transcript and rates four signals from 0 to 10: **Financial** (can they afford it), **Urgency** (do they need it now), **Authority** (can they decide), and **Engagement** (are they leaning in).

$$\text{Intent Index} = (\sum \text{signal} \times \text{weight} \div \sum \text{weights}) \div 2$$

Result: 0 to 5, one decimal.

WHY THIS EXISTS

Separating intent from score answers the question every sales manager asks: was this a bad call or a bad lead? A rep with a low close rate and consistently low intent has a marketing problem, not a coaching problem. This is also what makes the Marketing Intelligence feature possible.

The four weights default to 25% each and can be configured per organisation. The weights used are snapshotted onto each call at analysis time, so changing them later does not silently rewrite history.

OFTEN MISTAKEN FOR A BUG

The four signals are described **two different ways** in the product. The prompt that instructs the AI defines Financial as "does the prospect have budget available?", while the on-screen help text for the same signal asks "did the person react to the price with curiosity or resistance?". Budget availability and reaction to price overlap but are not the same thing. All four signals have this gap between the internal scoring definition and the customer-facing description. The customer guide follows the on-screen wording, since that is what users read — but the two should be reconciled so the AI is scoring exactly what the interface claims it scores.

KNOWN DEFECT — ALREADY REPORTED

The Intent Analysis dashboard **ignores the configured weights**. It fetches them from an endpoint that does not exist in production — the route was only ever built as a development mock — so the fetch fails, and the page silently falls back to equal weights of 25% each. An owner who sets Financial to 50% will see the setting save correctly and change nothing on the dashboard below it. The call detail screen is unaffected, because it reads the snapshot stored on the call itself.

11.2 When intent cannot be measured

OFTEN MISTAKEN FOR A BUG

If the AI fails to return an intent reading, the system assigns a value based purely on the outcome: closed → 4, partial → 3, not closed → 2, no outcome → 1. This is a fixed rule, not a measurement — the number looks like an assessment of the prospect but is really just a restatement of the result.

OFTEN MISTAKEN FOR A BUG

For calls analysed before this feature existed, the breakdown shown is a flat 5/5/5/5 placeholder. Those calls are excluded from the "Highest Priority Leads" table but **are** included in the radar chart averages, diluting them toward the middle with no visual marker.

11.3 The "Won" column

OFTEN MISTAKEN FOR A BUG

The Won column on the priority leads table reads the CRM's opportunity status, **not** the Stage 2 "paying" marker. Two different notions of "won" exist in the product and this table uses the automatic one. A call marked as paying by the owner will not necessarily show as won here.

12. Marketing Intelligence

Reads the best closed calls and suggests ad copy — headlines, angles, and objections to pre-empt.

Sample: 3 to 5 calls, chosen at random from the 10 highest-scoring closed calls. Refreshes weekly (cached 7 days).

WHY THIS EXISTS

Sampling only won calls, and only high-scoring ones, is deliberate: the goal is to extract the language that works when everything goes right, then feed it back into acquisition. It is not meant to be representative of all calls.

Confidence percentage **AI OPINION**

KNOWN DEFECT — ALREADY REPORTED

The confidence figure (0–100) shown on each suggestion is the AI grading its own answer — the prompt literally asks for "your honest read of how strong the signal is". Code only clamps it to a valid range; nothing verifies it against sales data. It is then rendered as a progress bar with High / Medium / Low labels, which reads as a statistical confidence interval. It is not one.

The "based on N calls" badge is real — it reflects the actual sample size.

13. Admin Panel and Billing

13.1 What the admin panel shows

Four cards: Total Organisations, Pending Approvals, Total Calls, and Average Score.

KNOWN DEFECT — ALREADY REPORTED

The Average Score card renders as a percentage — for example "4.2%". The value is converted to the 0–5 scale and then given a "%" suffix, which is neither the 84% it would be on the stored scale nor the "4.2/5" that was intended. Cosmetic, but visible on the first screen an admin sees.

13.2 MRR PLACEHOLDER

OFTEN MISTAKEN FOR A BUG

MRR is typed in by hand when an organisation is created, and can be overridden manually afterwards. It is not connected to Stripe, and it is **never summed anywhere** — no dashboard adds it up. The billing model moved to per-minute consumption and MRR was left behind as a loose field. If someone asks "where do I see total MRR?", the answer is that no such screen exists.

13.3 Customer Health PLACEHOLDER

KNOWN DEFECT — ALREADY REPORTED

Every organisation is created as "healthy" and nothing ever changes it. There is no rule, no threshold, no job — the column is displayed in the admin table but is decorative. A churning client will show as healthy indefinitely.

13.4 Billing — how clients are charged CALCULATED

billable minutes = call duration rounded UP to the next whole minute
calls under 30 seconds are not billed at all
cost = billable minutes × the organisation's rate

Default rate: \$0.0667 per minute

WHY THIS EXISTS

Rounding up per call rather than summing raw seconds is standard telecom-style billing and is simple to explain on an invoice. The 30-second floor exists so that misdials and instant hang-ups — which cost the platform almost nothing — do not appear as charges and trigger support tickets.

The rate is configurable per organisation, which supports negotiated deals.

KNOWN DEFECT — ALREADY REPORTED

A second, contradictory rate exists. The organisations grid on the admin panel calculates cost at **\$2.00 per minute** — a leftover from the original pricing model — while the billing screens use the real \$0.0667. The file holding the \$2 value even describes itself as "the single source of truth". The same organisation, in the same month, shows costs roughly 30 times apart depending on which admin screen you open.

13.5 LLM cost CALCULATED

$$\text{cost} = (\text{input tokens} \times \text{input price} + \text{output tokens} \times \text{output price}) \div 1,000,000$$

Prices live in a versioned database table. Editing a price never overwrites the old one — it deactivates it and inserts a new row with an effective date, so historical costs stay accurate. Current active pricing covers OpenAI (GPT-4o, GPT-4o-mini, GPT-4-turbo, GPT-4, GPT-3.5-turbo, Whisper) and Gemini (2.5 Pro, 2.5 Flash, 2.5 Flash-Lite, 2.0 Flash, 2.0 Flash-Lite).

WHY THIS EXISTS

This replaced an earlier estimate that simply assumed cost was 30% of revenue. Margin figures in the admin panel are now real, measured per API call.

Owners never see this. It is admin-only, filtered server-side.

13.6 Plan limits CALCULATED

Plan	Sales people	Calls per month	Key features
Starter	5	200	Manual upload, AI analysis, coaching email
Pro	15	1,000	Adds CRM webhook, automatic ingestion
Pro + RAG	Unlimited	Unlimited	Adds knowledge base, vector search

WHY THIS EXISTS

These limits are enforced by the database itself, not by the application, using locks that make the check atomic. That means two simultaneous invitations cannot both slip past the seat limit — a race condition that application-level checks would not catch. Owners do not count against the seat limit; only reps do.

The monthly call limit counts from the first of the calendar month, in UTC. There is no closing job — the month simply rolls over.

14. Known Defects — Do Not Report These

Every item below has been verified against the code and, where possible, against real production data. They are already identified. Reporting them again costs everyone time — but if you see behaviour that is *not* on this list and not explained elsewhere in this manual, that is worth raising.

#	Where	What is wrong	Impact
1	Scoring engine	Section weights and "critical" flags are saved but never affect the overall score, which is a plain average	High — script configuration silently does nothing
2	Dashboard → Correlation Engine	Calculates no correlation; shows average score with correlation labels	High — can point coaching in the wrong direction
3	Dashboard → Revenue Leak insight	Compares 0–100 scores against a 3.5 threshold meant for the 0–5 scale; prints "92/5"	High — card is statistically meaningless
4	Admin → organisations grid	Costs calculated at \$2.00/min instead of the real \$0.0667/min	High — two admin screens disagree by ~30x
5	Intent Analysis	Dashboard ignores configured weights; always uses 25/25/25/25 because the endpoint does not exist in production	High — a settings screen that does nothing
6	Script Intelligence	Health score, section scores and "+18%" uplifts are AI opinion presented as measurement	High — drives a real script-activation decision
7	Admin panel	Customer Health is always "healthy"; no rule ever changes it	Medium — decorative field
8	Admin panel	Average Score renders as "4.2%" — 0–5 value with a percent sign	Low — cosmetic
9	Call detail	Coaching notes textarea does not save	Medium — looks functional, silently loses input
10	Team Command Center	Scores recalculated live but deltas still read from a stale cache	Medium — correct score beside a wrong delta
11	Coaching recommendations	On AI failure, silently shows advice from a fictional demo persona	Medium — no indication it is not personalised
12	Marketing Intelligence	Confidence % is the AI grading itself, rendered as a statistical bar	Medium — false precision
13	Critical section alerts	Screen alerts below 40, coaching email alerts below 60	Low — email over-alerts relative to the app
14	Analytics → Rising Star	Compares organisation-wide first/last calls, credits the most recent rep	Low — demo-grade heuristic
15	Stage 2 (Actual Close)	Marker shows its own value, but no report consumes it and the intent-at-close snapshot is never read	Low — incomplete, not broken

16	Buying Intent	Each signal is defined one way in the AI prompt and described another way in the on-screen help (e.g. Financial: "has budget?" vs "reacted to price?")	Low — the AI may not score exactly what the UI claims
----	---------------	--	---

A note on scale drift

Several of the defects above share one root cause. The product originally scored calls on 0–5 and later migrated to 0–100. During that migration, a temporary rule was added in several places: *"if the value is 5 or less, assume it is on the old scale and multiply by 20"*. The migration finished; the rule was never removed.

Production data confirms this is not theoretical: **4.9% of real section scores are 5 or below** on the 0–100 scale — genuinely terrible sections. Every one of them is being multiplied by 20, turning a catastrophic 3/100 into a comfortable 60/100. The rule inverts precisely the worst calls, which are the ones a coach most needs to see.

15. Quick Reference

The same metric, calculated differently in different places

This is the most common cause of "these two screens disagree" reports. In most cases both numbers are correct — they answer different questions.

Close rate shown on...	Covers
Dashboard headline card	Whole organisation, all time, call by call
Dashboard trend / delta	Last 6 weeks maximum, weekly buckets
Team Health rows	Rep's lifetime, from cache; delta = last 7 days vs. everything before
Rep's My Page	Selected window of 2, 4 or 6 weeks, volume-weighted
AI coaching recommendations	Rep's last 20 calls
Analytics leaderboard	All loaded calls, grouped by rep name

Sample sizes behind each AI feature

Feature	Calls analysed	Refresh
Script Gap Detection	3 (full transcripts)	Every 7 days
Script Intelligence	Up to 7 (truncated to 1,500 chars)	On demand
Marketing Intelligence	3–5, random from top 10 closed	Every 7 days
Coaching recommendations	Last 20 per rep	On demand
Dashboard rubric & trends	200 most recent	Every page load
Analytics	~1,000 most recent	Every page load

Key numbers at a glance

Rule	Value
Internal score scale	0–100
Displayed score scale	0–5 (stored ÷ 20)
Green / amber / red	≥85 / 70–84 / <70
Perfect call	≥95
Overall score	Simple average of sections
Billing rate (default)	\$0.0667 per minute
Minimum billable call	30 seconds
Billing rounding	Up, to the whole minute
Intent Index	0–5, from 4 signals scored 0–10
Default intent weights	25% each

Deprecated routes

The `/overview` screen has been removed. The owner's home is `/dashboard`. Old links to `/overview` redirect there automatically rather than erroring, so bookmarks keep working.